

Pre requisite: Install k8s cluster (minikube) + kubectl

Notes: minikube can be deployed as a VM, a container

Start it using minikube start --driver=docker OR minikube start --driver=virtualbox

This makes kubectl configured to use "minikube" cluster and "default" namespace my default

- 1- Create a pod with the name "imperative-nginx" and with the image `nginx` and `latest` tag. using Imperative command (not yaml).
- 2- Create a pod with the name `webserver` and with the image "nginx123" Use a pod-definition YAML file.
- 3- What is the `nginx` pod status?
- 4- Change the `nginx` pod image to "nginx" check the status again
- 5- How many pods are running in the system? Type the command to show this
- 6- What does READY column in the output of `get pods` command indicate?
- 7- Delete first pod named `imperative-nginx` you just created. Type the command to do this
- 8- Which node is pod named `webserver` running on (list two commands to do this)
- 9- Get a shell to the running container i.e ssh into it (figure out the command)
- 10- Run `cat /etc/os-release` inside the container
- 11- Exit from the shell (`/bin/bash`) session
- 12- Get logs of pod, what are logs and what they are used for?
- 13- How many ReplicaSets exist on the system?
- 14- create a ReplicaSet
withname= `replica-set-1`
image= `busybox`
replicas= 3
- 15- Scale the ReplicaSet `replica-set-1` to 5 PODs.
- 16- How many PODs are READY in the `replica-set-1`?
- 17- Delete any one of the 5 PODs then check How many PODs exist now?
Why are there still 5 PODs, even after you deleted one?